

Table of Contents

- [Library Import & Settings](#)
- [Load Data into Pandas](#)
- [Basic Properties of the Dataset](#)
- [Exploring Column Summaries](#)
 - [Summary for Categorical Features](#)
 - [Summary for Numerical Features](#)
- [Exploring Text Categories](#)
 - [Visualizing Frequency Distributions](#)
- [Exploring Word Frequencies](#)
 - [Creating a List of Tokens from a List of Documents](#)
 - [Counting Frequencies with a Counter](#)
 - [Using Word Clouds](#)
- [Exploring Text Complexity](#)

Modern Text Mining with Python (Part 2 of 5): Data Exploration with Pandas

This notebook contains the code examples for the blog post on [Medium.com](#).

To get it running, you will need the SQLite Database `selfposts.db`, which was created in the first notebook of this series.

Library Import & Settings

```
In [ ]: # filter warnings on depreciation etc.  
import warnings  
warnings.filterwarnings("ignore")
```

```
In [ ]: # import pandas, numpy  
import pandas as pd  
import numpy as np  
  
# adjust pandas display  
pd.options.display.max_columns = 30  
pd.options.display.max_rows = 100  
pd.options.display.float_format = '{:.2f}'.format  
pd.options.display.precision = 2  
pd.options.display.max_colwidth = -1
```

```
In [ ]: # Import matplotlib and seaborn and adjust some defaults
%matplotlib inline
%config InlineBackend.figure_format = 'svg'

from matplotlib import pyplot as plt
plt.rcParams['figure.dpi'] = 100

import seaborn as sns
sns.set_style("whitegrid")
```

Load Data into Pandas

```
In [ ]: import sqlite3

sql = """
SELECT p.*,
       c.category_1 as category, c.category_2 as subcategory
FROM posts_nlp p, categories c
WHERE p.subreddit = c.subreddit
"""

with sqlite3.connect('/home/deepcode/Downloads/selfposts.db') as con:
    df = pd.read_sql_query(sql, con)
```

Basic Properties of the Dataset

```
In [ ]: # list column names and datatypes
df.dtypes
```

```
Out[ ]: index          int64
id              object
subreddit       object
title           object
selftext        object
selftext_clean  object
selftext_lemma  object
selftext_nouns  object
selftext_adjectives object
selftext_verbs  object
selftext_nav    object
no_tokens       float64
category        object
subcategory      object
dtype: object
```

```
In [ ]: # select a sample of some data frame columns
df[['id', 'subreddit', 'title', 'selftext_clean']] \
    .sample(2, random_state=42)
```

Out[]:

	id	subreddit	title	selftext_clean
1501	4rwc67	AskLiteraryStudies	ideas for a tragedy survey	Hello folks! I am due to be teaching an undergrad (no pre-reqs) survey of European tragedy next semester, and I'm looking for some ideas of primary texts to put on the syllabus. The course has a pretty set component in the form of classical texts, but the rest of it is freer. I would especially welcome any suggestions or tips on 1) renaissance/baroque plays, 2) something from Calderon, 3) German classicism, 4) modern drama. If anyone can think of plays (the definition of "tragedy" can be fudged) by women or ethnic writers, that would be really helpful. This is the first time I'm teaching something like this, and it's been dropped in my lap rather suddenly, so I would be very grateful for suggestions from the more experienced teachers among you! Thanks in advance!
2586	7gdvsv	breastfeeding	Question about a very hungry child	Hey guys, My kid is 9 months old and still wakes up every 2 hours throughout the night to breastfeed. At this point, I co-sleep with her to get some sleep through the night and now I really like the arrangement. Is there a way or technique I could use where I could continue to co-sleep but the baby doesn't feed that much at night? She is 10% weight and is EBF. She also started solids at 6 months and eats three meals a day at this point.

In []: *# length of a dataframe*
len(df)

Out[]: 5000

In []: *# number of values per column*
df.count()

Out[]:

index	5000
id	5000
subreddit	5000
title	5000
selftext	5000
selftext_clean	5000
selftext_lemma	5000
selftext_nouns	5000
selftext_adjectives	5000
selftext_verbs	5000
selftext_nav	5000
no_tokens	5000
category	5000
subcategory	5000
dtype: int64	

In []: *# size info, including memory consumption*
df.info(memory_usage='deep')

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 14 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   index                 5000 non-null   int64
 1   id                    5000 non-null   object
 2   subreddit             5000 non-null   object
 3   title                 5000 non-null   object
 4   selftext              5000 non-null   object
 5   selftext_clean        5000 non-null   object
 6   selftext_lemma        5000 non-null   object
 7   selftext_nouns        5000 non-null   object
 8   selftext_adjectives   5000 non-null   object
 9   selftext_verbs        5000 non-null   object
10   selftext_nav          5000 non-null   object
11   no_tokens             5000 non-null   float64
12   category              5000 non-null   object
13   subcategory           5000 non-null   object
dtypes: float64(1), int64(1), object(12)
memory usage: 20.0 MB

```

Exploring Column Summaries

The pandas `describe` method computes statistical summaries for each of the columns of a dataframe. The results are different for categorical and numerical features.

Summary for Categorical Features

```
In [ ]: columns = [col for col in df.columns if not col.startswith('self')]
columns
```

```
Out[ ]: ['index', 'id', 'subreddit', 'title', 'no_tokens', 'category', 'subcategory']
```

```
In [ ]: # describe categorical columns of type np.object
df[['category', 'subcategory', 'subreddit']] \
    .describe(include=object) \
    .transpose()
```

```
Out[ ]:
```

	count	unique	top	freq
category	5000	39	video_game	488
subcategory	5000	1003	logo	15
subreddit	5000	1003	logorequests	15

```
In [ ]: df['subreddit'].value_counts()[:10]
```

```
Out[ ]: logorequests      15
        samuraijack      12
        RocketLeague     12
        Tinder           12
        backpacking      12
        Tekken           12
        solar            11
        gigantic         11
        DungeonWorld     11
        speedrun         11
        Name: subreddit, dtype: int64
```

Summary for Numerical Features

```
In [ ]: # describe numerical columns
        df.describe().transpose()
```

```
Out[ ]:
```

	count	mean	std	min	25%	50%	75%	max
index	5000.00	2499.50	1443.52	0.00	1249.75	2499.50	3749.25	4999.00
no_tokens	5000.00	163.90	132.00	20.00	79.00	117.00	193.00	943.00

Exploring Text Categories

```
In [ ]: # number of unique values = count distinct
        df['category'].nunique()
```

```
Out[ ]: 39
```

```
In [ ]: # group by category, count distinct subreddits and posts
        cat_df = df.groupby('category') \
                .agg({'subreddit': pd.Series.nunique,
                     'id': pd.Series.count}) \
                .rename(columns={'subreddit': 'num_subreddits',
                                'id': 'num_posts'}) \
                .sort_values('num_subreddits', ascending=False)

        # show top 5 records
        cat_df.head(5)
```

```
Out[ ]:
```

	num_subreddits	num_posts
category		
video_game	97	488
tv_show	68	323
health	57	290
profession	56	257
software	51	244


```
In [ ]: cat_df.describe()
```

```
Out[ ]:
```


	num_subreddits	num_posts
count	39.00	39.00
mean	25.72	128.21
std	19.45	94.74
min	5.00	30.00
25%	13.00	71.50
50%	20.00	95.00
75%	29.50	152.00
max	97.00	488.00

Visualizing Frequency Distributions

```
In [ ]: # horizontal boxplot of a dataframe column  
cat_df[['num_posts']].plot(kind='box', vert=False, figsize=(6, 2));
```



```
In [ ]: # bar chart of a dataframe column  
cat_df[['num_subreddits']].plot(kind='bar', figsize=(7, 4));
```



Exploring Word Frequencies

```
In [ ]: # create a data frame slice
sub_df = df[df['subreddit']=='TheSimpsons']

# sample cleaned text and tokens tagged as nouns
sub_df[['selftext_clean', 'selftext_nouns']].sample(2)
```

		selftext_clean	selftext_nouns
3098	Just saw on Youtube as a 3 min clip: Homer finds a antique green car in his garage (was green) Marge was upset that homer didn't do his chores so he did it on "laminated paper" he took car everywhere, that is where the episode had ended and could not find again help please		Youtube min clip Homer car garage Marge homer chore paper car episode
823	Okay, I've been researching for two hours and I can't find this episode that I remember. It pops in my head every once in a while and I've never been able to find it, you guys are my last hope. I'm pretty sure it was within the last... 6 years? maybe more. All that I remember is Lisa meets a new, possibly weird/creepy boy who I thiink has been living at the school? He may have lost his parents. I'm pretty sure I remember him being into creepy stuff but it being a full episode and not a treehouse of horror or anything. If anyone has any idea what I'm talking about I'd be eternally grateful, I have no idea why but this one particular episode eludes me. (I also swear I remember him being voiced by Elijah Wood but apparently he's never been a guest on the show ヽ(ヾ)ノ		hour episode head while guy hope year Lisa boy school parent stuff episode treehouse horror idea idea episode Elijah Wood guest show ヽヾ

Creating a List of Tokens from a List of Documents

```
In [ ]: def my_tokenizer(text):
        return text.split() if text != None else []
```

```
In [ ]: # transform list of documents into a single list of tokens
tokens = sub_df.selftext_nouns.map(my_tokenizer).sum()
```

```
In [ ]: print(tokens[:200])
```

```
['episode', 'time', 'other', 'instance', 'proposal', 'Putlocker', 'week',
'tv', 'Bart', 'Thanksgiving', 'episode', 'other', 'show', 'girlfriend', 'ki
d', 'point', 'go', 'show', 'movie?there', 'thing', 'reference', 'continuit
y', 'quality', 'reference', 'back', 'movie', 'season', 'Mr.', 'Burns', 'jok
e', 'consequence', 'callback', 'order', 'problem', 'quality', 'matching',
'movie', 'lot', 'episode', 'show', 'point', 'season', 'dip', 'movie', 'en
d', 'cap', 'movie', 'cartooniness', 'vibe', 'show', 'hour', 'episode', 'hea
d', 'while', 'guy', 'hope', 'year', 'Lisa', 'boy', 'school', 'parent', 'stu
ff', 'episode', 'treehouse', 'horror', 'idea', 'idea', 'episode', 'Elijah',
'Wood', 'guest', 'show', '``\\', 'ツ', 'post', 'time', 'episode', 'episode',
'Homer', 'City', 'New', 'York', 'season', 'episode', 'Youtube', 'min', 'cli
p', 'Homer', 'car', 'garage', 'Marge', 'homer', 'chore', 'paper', 'car', 'e
pisode', 'name', 'episode', 'Homer', 'bully', 'Kearney', 'joke', 'Homer',
'line', 'guy', 'thank', 'episode', 'Simpsons', 'episode', 'Lisa', 'Gaga',
'braindead', 'STUPID', 'paddling', 'episode', 'scene', 'Homer', 'Gill', 'ca
r', 'dog', 'way', 'show', 'minute', 'episode', 'court', 'minute', 'court',
'thing', 'Homer', 'dog', 'Springfield', 'dog', 'treehouse', 'horror', 'do
g', 'control', 'city', 'town', 'Gill', 'dog', 'Marge', 'Chihuahua', 'Deus',
'Ex', 'Machina', 'Marge', 'Chihuahua', 'death', 'moronic', 'joke', 'episod
e', 'reference', 'Turbo', 'WHY', 'REALLY', 'ASK', 'fucking', 'episode', 'tu
rbo', 'joke', 'month', 'production', 'time', 'hell', 'Turbo', 'month', 'tur
bo', 'episode', 'Treehouse', 'Horror', 'breakneck', 'pace', 'Treehouse', 'H
orror', 'heart', 'episode', 'GOD']
```

Counting Frequencies with a Counter

```
In [ ]: from collections import Counter
```

```
counter = Counter(tokens)
counter.most_common(20)
```



```
Out[ ]: [('episode', 19),
        ('show', 6),
        ('Homer', 6),
        ('dog', 5),
        ('movie', 4),
        ('joke', 4),
        ('time', 3),
        ('reference', 3),
        ('season', 3),
        ('car', 3),
        ('Marge', 3),
        ('other', 2),
        ('point', 2),
        ('thing', 2),
        ('quality', 2),
        ('guy', 2),
        ('Lisa', 2),
        ('treehouse', 2),
        ('horror', 2),
        ('idea', 2)]
```

```
In [ ]: df.category.unique()
```

```
Out[ ]: array(['writing/stories', 'tv_show', 'autos', 'hardware/tools',
              'electronics', 'video_game', 'crypto', 'sports', 'hobby',
              'appearance', 'card_game', 'drugs', 'advice/question',
              'social_group', 'anime/manga', 'sex/relationships', 'software',
              'health', 'other', 'animals', 'arts', 'programming', 'rpg',
              'books', 'parenting', 'education', 'company/website', 'profession',
              'music', 'politics/viewpoint', 'stem', 'travel', 'geo',
              'religion/supernatural', 'board_game', 'movies', 'food/drink',
              'finance/money', 'meta'], dtype=object)
```

```
In [ ]: print([t[0] for t in counter.most_common(200)])
```

```
['episode', 'show', 'Homer', 'dog', 'movie', 'joke', 'time', 'reference',
'season', 'car', 'Marge', 'other', 'point', 'thing', 'quality', 'guy', 'Lisa',
'treehouse', 'horror', 'idea', 'Gill', 'minute', 'court', 'Chihuahua',
'Turbo', 'turbo', 'month', 'Treehouse', 'Horror', 'instance', 'proposal',
'Putlocker', 'week', 'tv', 'Bart', 'Thanksgiving', 'girlfriend', 'kid', 'go',
'movie?there', 'continuity', 'back', 'Mr.', 'Burns', 'consequence', 'callback',
'order', 'problem', 'matching', 'lot', 'dip', 'end', 'cap', 'cartooniness',
'vibe', 'hour', 'head', 'while', 'hope', 'year', 'boy', 'school',
'parent', 'stuff', 'Elijah', 'Wood', 'guest', '``', 'ツ', 'post', 'City',
'New', 'York', 'Youtube', 'min', 'clip', 'garage', 'homer', 'chore', 'paper',
'name', 'bully', 'Kearney', 'line', 'thank', 'Simpsons', 'Gaga', 'braindead',
'STUPID', 'paddling', 'scene', 'way', 'Springfield', 'control', 'city',
'town', 'Deus', 'Ex', 'Machina', 'death', 'moronic', 'WHY', 'REALLY',
'ASK', 'fucking', 'production', 'hell', 'breakneck', 'pace', 'heart', 'GOD']
```

```
In [ ]: from spacy.lang.en.stop_words import STOP_WORDS


def remove_stopwords(tokens):
    """Remove stopwords from a list of tokens."""
    return [t for t in tokens if t not in STOP_WORDS]
```

```
# rebuild counter
counter = Counter(remove_stopwords(tokens))
```

```
2023-03-19 13:51:03.468638: I tensorflow/core/platform/cpu_feature_guard.cc:193] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: AVX2 FMA
To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.
2023-03-19 13:51:03.635675: W tensorflow/compiler/xla/stream_executor/platform/default/dso_loader.cc:64] Could not load dynamic library 'libcudart.so.11.0'; dlerror: libcudart.so.11.0: cannot open shared object file: No such file or directory
2023-03-19 13:51:03.635710: I tensorflow/compiler/xla/stream_executor/cuda/cudart_stub.cc:29] Ignore above cudart dlerror if you do not have a GPU set up on your machine.
2023-03-19 13:51:04.316902: W tensorflow/compiler/xla/stream_executor/platform/default/dso_loader.cc:64] Could not load dynamic library 'libnvinfer.so.7'; dlerror: libnvinfer.so.7: cannot open shared object file: No such file or directory
2023-03-19 13:51:04.316998: W tensorflow/compiler/xla/stream_executor/platform/default/dso_loader.cc:64] Could not load dynamic library 'libnvinfer_plugin.so.7'; dlerror: libnvinfer_plugin.so.7: cannot open shared object file: No such file or directory
2023-03-19 13:51:04.317004: W tensorflow/compiler/tf2tensorrt/utils/py_utils.cc:38] TF-TRT Warning: Cannot dlopen some TensorRT libraries. If you would like to use Nvidia GPU with TensorRT, please make sure the missing libraries mentioned above are installed properly.
2023-03-19 13:51:05.057663: W tensorflow/compiler/xla/stream_executor/platform/default/dso_loader.cc:64] Could not load dynamic library 'libcuda.so.1'; dlerror: libcuda.so.1: cannot open shared object file: No such file or directory
2023-03-19 13:51:05.057688: W tensorflow/compiler/xla/stream_executor/cuda/cuda_driver.cc:265] failed call to cuInit: UNKNOWN ERROR (303)
2023-03-19 13:51:05.057712: I tensorflow/compiler/xla/stream_executor/cuda/cuda_diagnostics.cc:156] kernel driver does not appear to be running on this host (dc-central): /proc/driver/nvidia/version does not exist
```

```
In [ ]: # convert list of tuples into data frame
freq_df = pd.DataFrame.from_records(counter.most_common(20),
                                     columns=['token', 'count'])

# create bar plot
freq_df.plot(kind='bar', x='token');
```



Using Word Clouds

```
In [ ]: %matplotlib inline
import matplotlib.pyplot as plt
```

```
In [ ]: %%capture
%pip install wordcloud
```

```
In [ ]: from wordcloud import WordCloud

def wordcloud(counter):
    """A small wordcloud wrapper"""
    wc = WordCloud(width=1200, height=800,
                   background_color="white",
                   max_words=200)
    wc.generate_from_frequencies(counter)

    # Plot
    fig=plt.figure(figsize=(6, 4))
    plt.imshow(wc, interpolation='bilinear')
    plt.axis("off")
    plt.tight_layout(pad=0)
    plt.show()
```

```
In [ ]: # create wordcloud
```

```
wordcloud(counter)
```

Comparing to Sopranos


```
In [ ]: tokens2 = df[df['subreddit']=='thesopranos'].selftext_nouns \
        .map(my_tokenizer).sum()

counter2 = Counter(remove_stopwords(tokens2))
wordcloud(counter2)
```

Exploring Text Complexity

```
In [ ]: df['no_tokens'] = df.selftext_lemma\  
        .map(lambda l: 0 if l==None else len(l.split()))
```

```
In [ ]: # mean number of tokens by category  
df.groupby(['category']) \  
  .agg({'no_tokens': 'mean'}) \  
  .sort_values(by='no_tokens', ascending=False) \  
  .plot(kind='bar', figsize=(7,4));
```



```
In [ ]: # render plots as retina or png, because svg is very slow
%config InlineBackend.figure_format = 'retina'

import seaborn as sns

def multi_boxplot(data, x, y, ylim = None):
    '''Wrapper for sns boxplot with cut-off functionality'''
    # plt.figure(figsize=(30, 5))
    fig, ax = plt.subplots()
    plt.xticks(rotation=90)


    # order boxplots by median
    ordered_values = data.groupby(x)[[y]] \
        .median() \
        .sort_values(y, ascending=False) \
        .index

    sns.boxplot(x=x, y=y, data=data, palette='Set2',
                order=ordered_values)

    fig.set_size_inches(11, 6)

    # cut-off y-axis at value ylim
    ax.set_ylim(0, ylim)
```


```
In [ ]: multi_boxplot(df, 'category', 'no_tokens');
```



```
In [ ]: # print text of outliers
df['selftext_lemma'][df.no_tokens > 1500]
```

```
Out[ ]: Series([], Name: selftext_lemma, dtype: object)
```

```
In [ ]: # cut-off diagram at y=500
multi_boxplot(df, 'category', 'no_tokens', ylim=500)
```



```
In [ ]: # comparing subreddits within a single category
multi_boxplot(df[df.category=='sex/relationships'],
              'subreddit', 'no_tokens', ylim=700)
```

